Assignment 3:

Provision VPC in us-east-1 region with 2 public and private subnets on two availability zones. Deploy Dynamic website presentation layer on ec2 instance on one of the public subnet. Spin up RDS (Managed Database) on private subnet and connect web layer to database layer using connection string. Finally deploy dynamic website.

Below is the summary of all steps for this assignment.

1. Create EC2 with all the required software to be installed. Check if the webserver is working with web page access.
2. Create new AMI with this instance.
3. Create VPC
    a. With Public subnet - web server associated with this → website is deployed here
    b. With Private subnet – Database server associated with this → DB server running
    c. Default route table
4. Create target group - To route traffic to the targets in a target group, specify the target group in an action when you create a listener
5. Create load balancer – public zone
6. Create a launch configuration and Auto scaling
    a. Create an Auto Scaling group using a launch configuration
    b. Note : Enable cloud watch, monitor health check

Once we finish all the above steps, we can access the website to see the load balance is working or not.

Also we can stop instance to see the request is going to available to webserver.

Below is the example image for the VPC with private and public subnet.

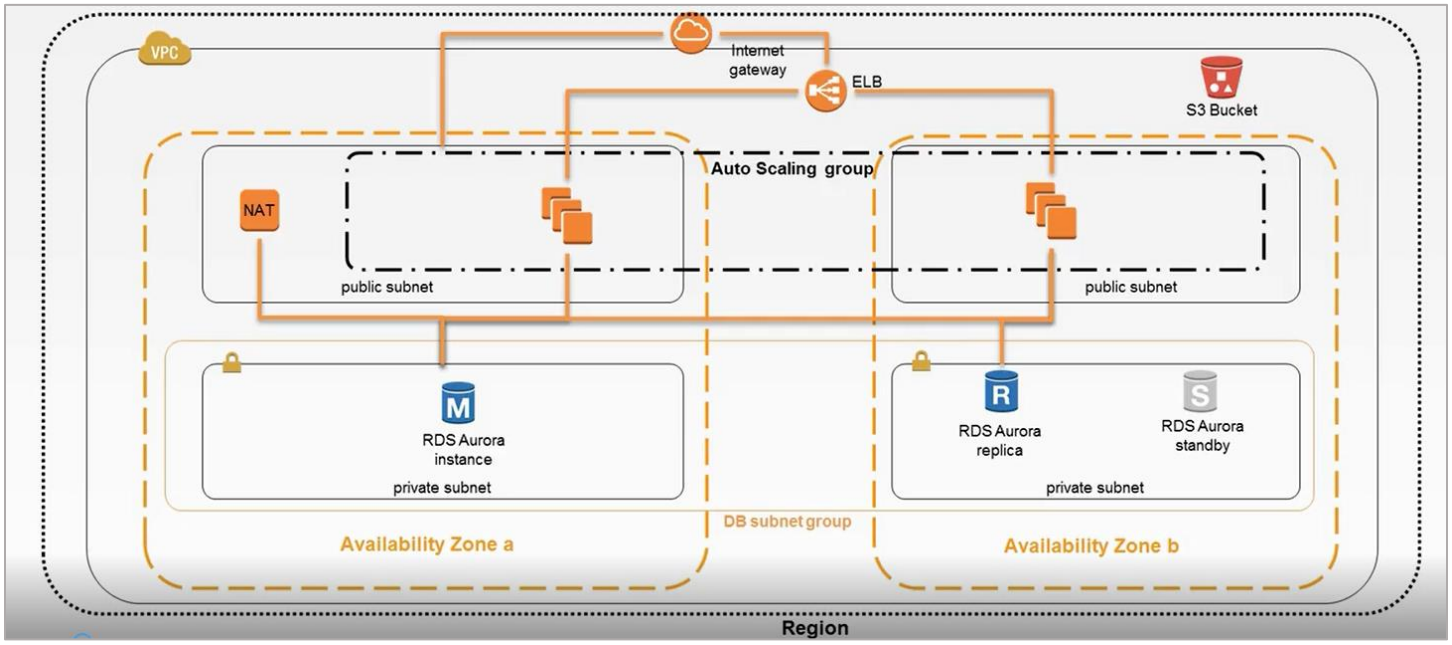

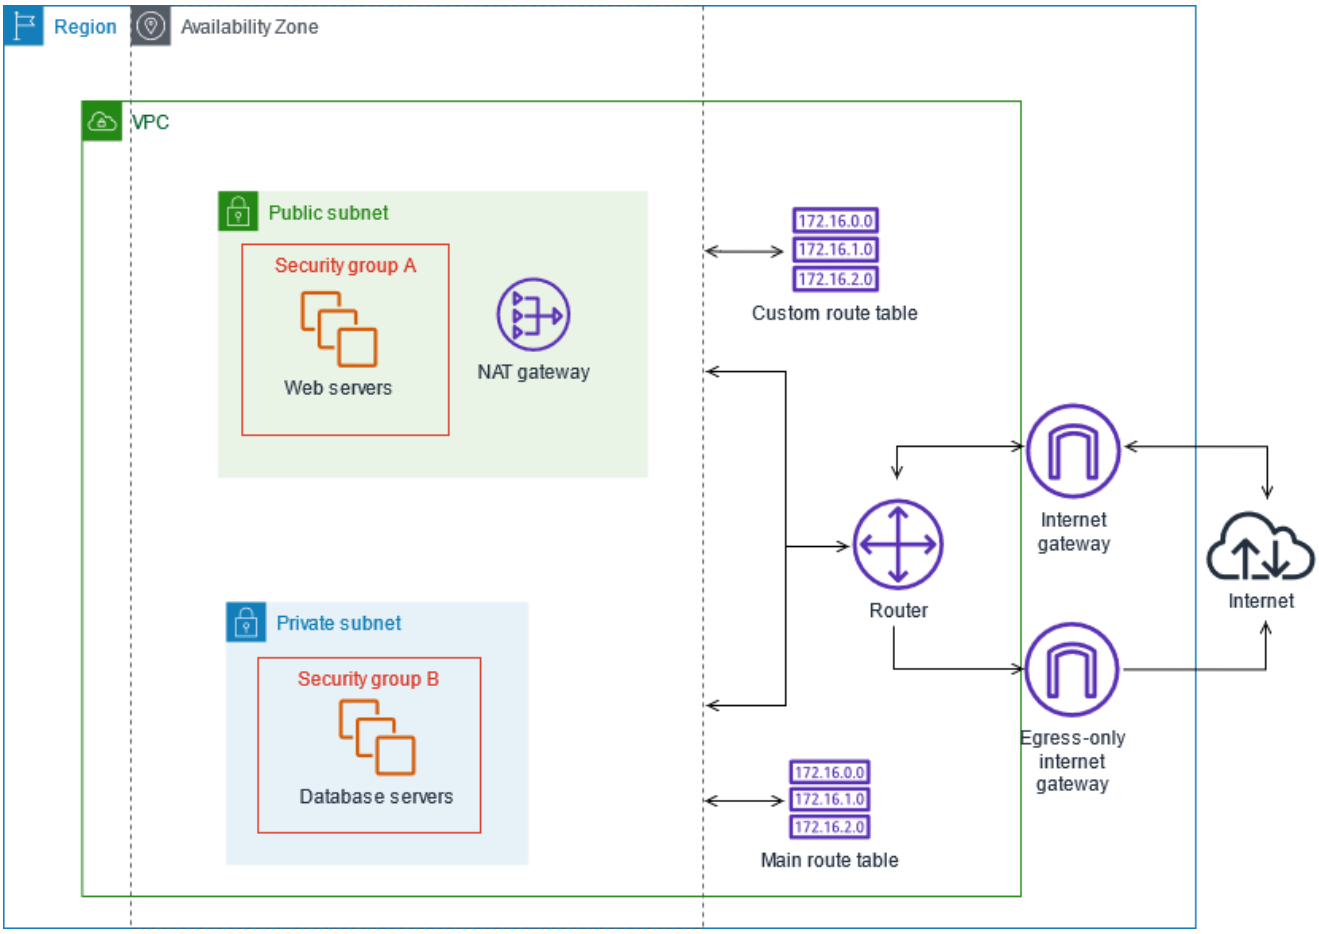
Region
Availability Zone
VPC
Public subnet
Security group A
Web servers
NAT gateway
172.16.0.0
172.16.1.0
172.16.2.0
Custom route table
Internet gateway
Router
Internet
Private subnet
Security group B
Database servers
Egress-only internet gateway
172.16.0.0
172.16.1.0
172.16.2.0
Main route table

1. Create VPC

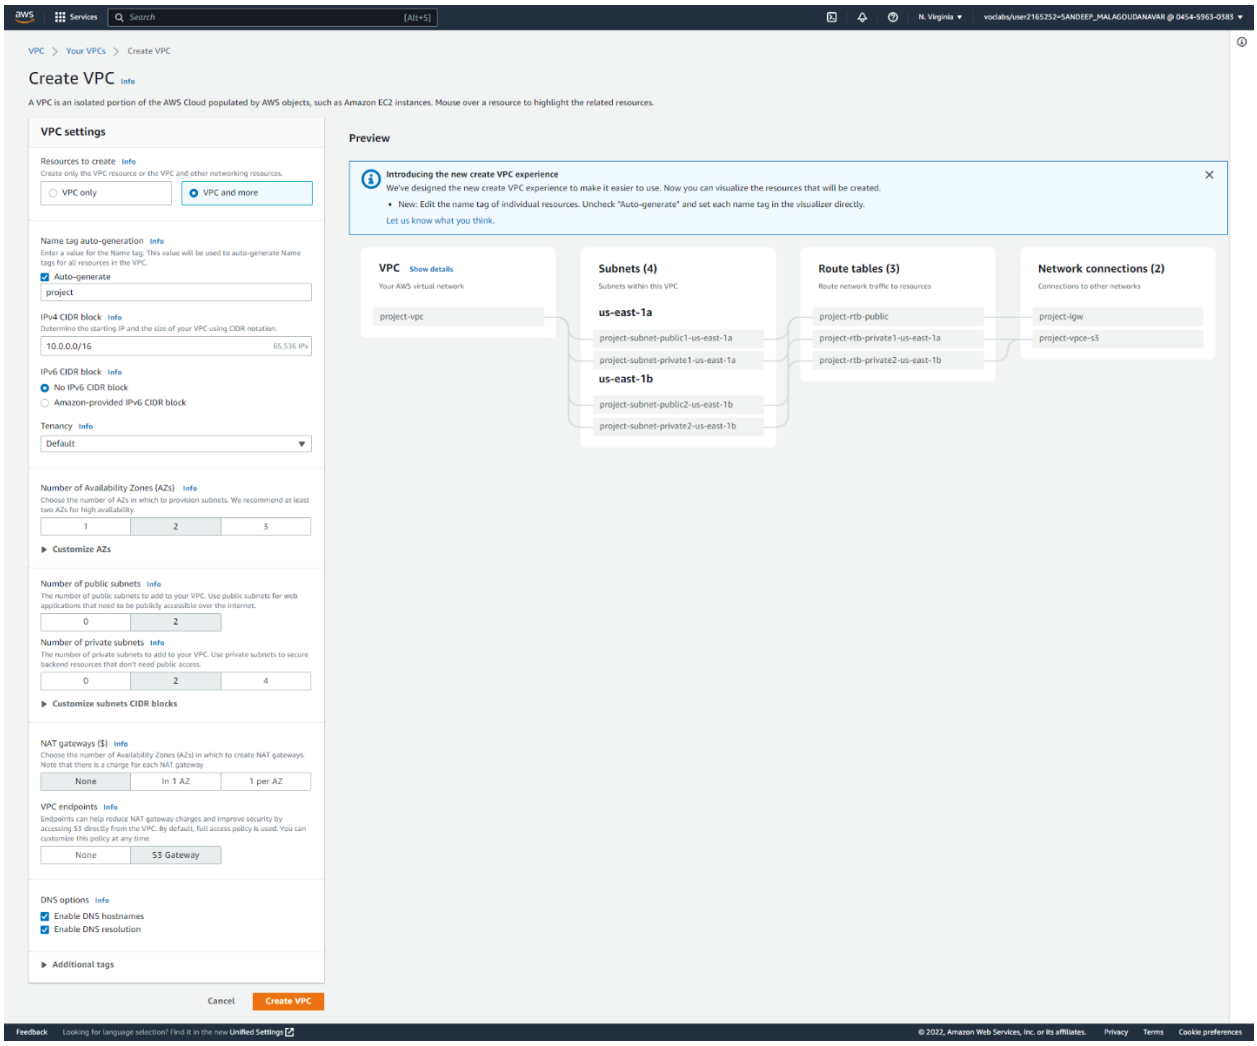
Services
Search
[Alt+S]
N. Virginia
VPC > Your VPCs > Create VPC
Create VPC Info
A VPC is an isolated portion of the AWS Cloud populated by AWS objects, such as Amazon EC2 instances. Mouse over a resource to highlight the related resources.
VPC settings
Resources to create Info
Create only the VPC resource or the VPC and other networking resources.
VPC only
VPC and more
Name tag auto-generation Info
Enter a value for the Name tag. This value will be used to auto-generate Name tags for all resources in the VPC.
Auto-generate
project
IPv4 CIDR block Info
Determine the starting IP and the size of your VPC using CIDR notation.
10.0.0.0/16
65,536 IPs
IPv6 CIDR block Info
No IPv6 CIDR block
Amazon-provided IPv6 CIDR block
Tenancy Info
Default
Number of Availability Zones (AZs) Info
Choose the number of AZs in which to provision subnets. We recommend at least two AZs for high availability.
1
2
3
Customize AZs
Number of public subnets Info
The number of public subnets to add to your VPC. Use public subnets for web applications that need to be publicly accessible over the internet.
0
2
Number of private subnets Info
The number of private subnets to add to your VPC. Use private subnets to secure backend resources that don't need public access.
0
2
4
Customize subnets CIDR blocks
NAT gateways ($) Info
Choose the number of Availability Zones (AZs) in which to create NAT gateways. Note that there is a charge for each NAT gateway.
None
In 1 AZ
1 per AZ
VPC endpoints Info
Endpoints can help reduce NAT gateway charges and improve security by accessing S3 directly from the VPC. By default, full access policy is used. You can customize this policy at any time.
None
S3 Gateway
DNS options Info
Enable DNS hostnames
Enable DNS resolution
Additional tags
Cancel
Create VPC
Preview
Introducing the new create VPC experience
We've designed the new create VPC experience to make it easier to use. Now you can visualize the resources that will be created.
New: Edit the name tag of individual resources. Uncheck "Auto-generate" and set each name tag in the visualizer directly.
Let us know what you think.
VPC Show details
Your AWS virtual network
project-vpc
Subnets (4)
Subnets within this VPC
us-east-1a
project-subnet-public1-us-east-1a
project-subnet-private1-us-east-1a
us-east-1b
project-subnet-public2-us-east-1b
project-subnet-private2-us-east-1b
Route tables (3)
Route network traffic to resources
project-rtb-public
project-rtb-private1-us-east-1a
project-rtb-private2-us-east-1b
Network connections (2)
Connections to other networks
project-igw
project-vpce-s3
Feedback
Looking for language selection? Find it in the new Unified Settings
© 2022, Amazon Web Services, Inc. or its affiliates.
Privacy
Terms
Cookie preferences

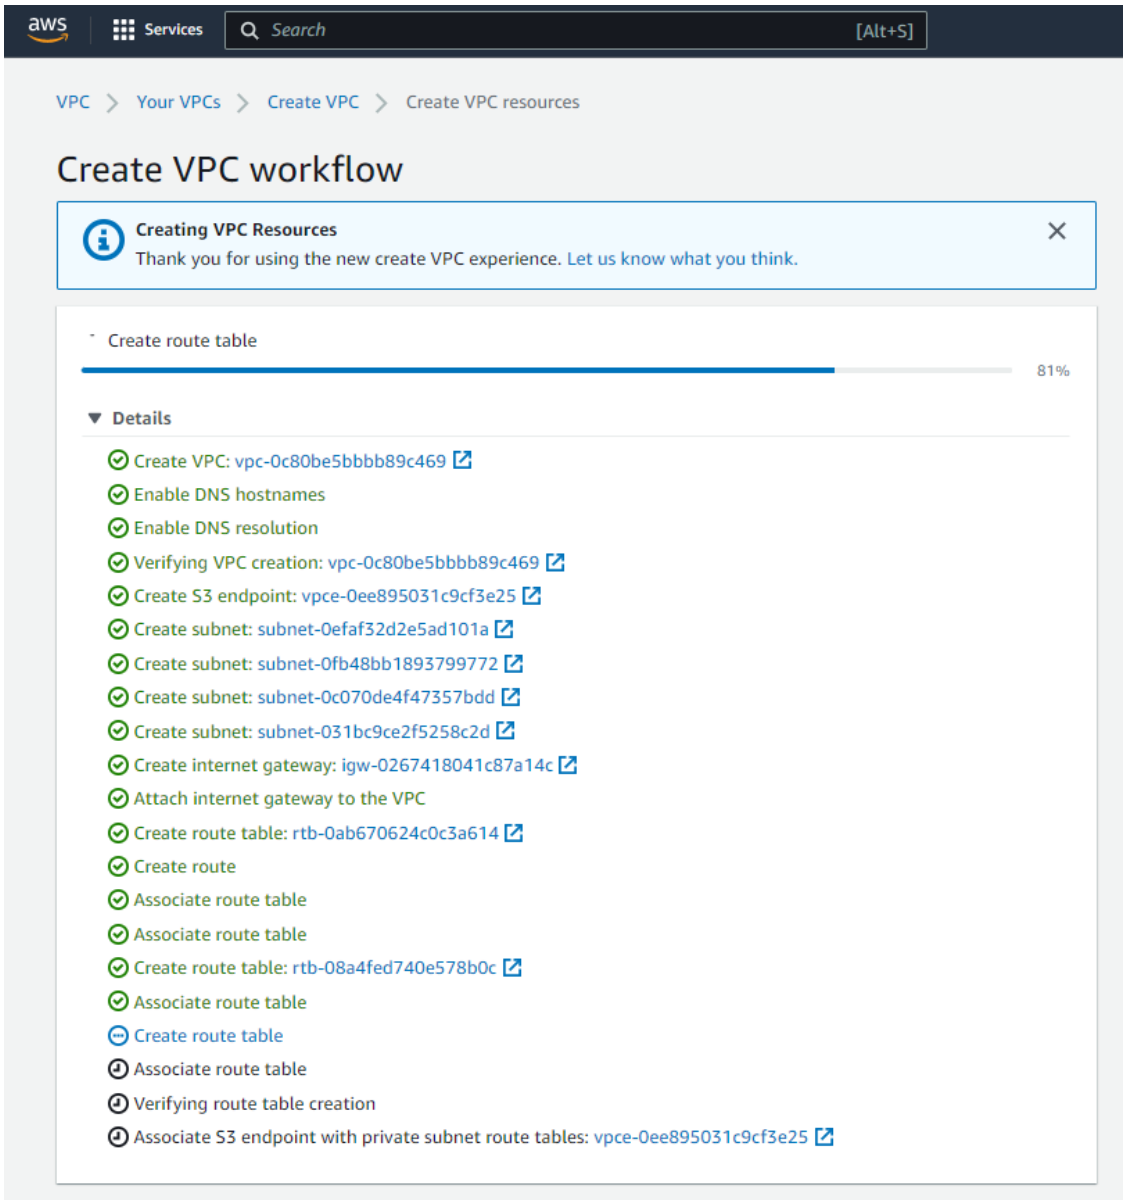

2. Create a target group

By default, the load balancer sends requests to registered targets using the port and protocol that you specified for the target group. To route traffic to the targets in a target group, specify the target group in an action when you create a listener or create a rule for your listener. Add or remove targets from your target group at any time.

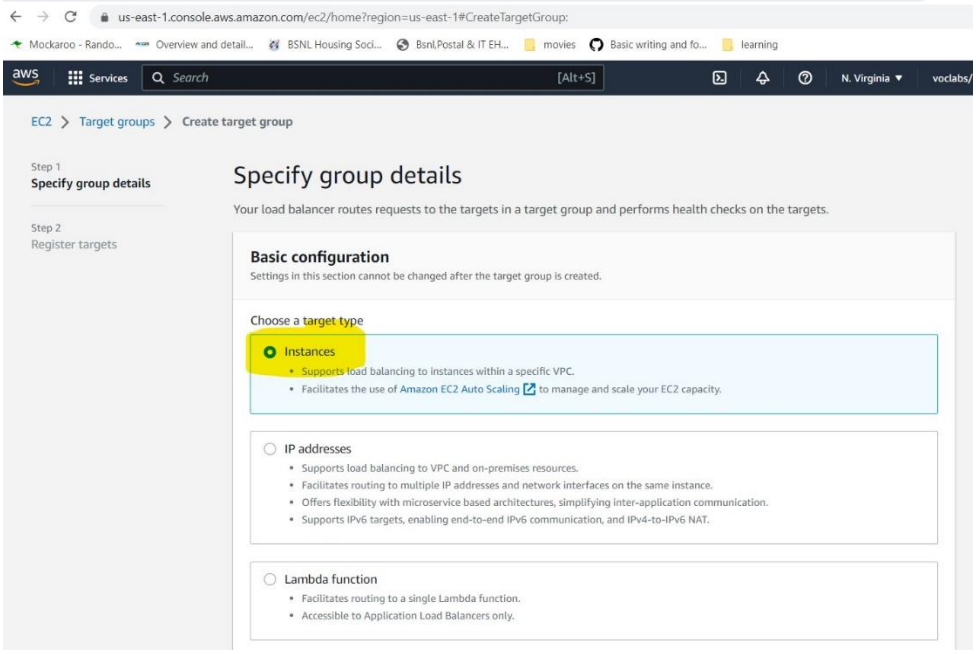
us-east-1.console.aws.amazon.com/ec2/home?region=us-east-1#CreateTargetGroup:
Mockaroo - Rando...
Overview and detail...
BSNL Housing Soci...
Bsnl,Postal & IT EH...
movies
Basic writing and fo...
learning
aws
Services
Search
[Alt+S]
N. Virginia
voclabs/
EC2 > Target groups > Create target group
Step 1
Specify group details
Step 2
Register targets
Specify group details
Your load balancer routes requests to the targets in a target group and performs health checks on the targets.
Basic configuration
Settings in this section cannot be changed after the target group is created.
Choose a target type
Instances
Supports load balancing to instances within a specific VPC.
Facilitates the use of Amazon EC2 Auto Scaling to manage and scale your EC2 capacity.
IP addresses
Supports load balancing to VPC and on-premises resources.
Facilitates routing to multiple IP addresses and network interfaces on the same instance.
Offers flexibility with microservice based architectures, simplifying inter-application communication.
Supports IPv6 targets, enabling end-to-end IPv6 communication, and IPv4-to-IPv6 NAT.
Lambda function
Facilitates routing to a single Lambda function.
Accessible to Application Load Balancers only.

[Alt+S]

target group

Specify group details

Your load balancer routes requests to the targets in a target group and performs health checks on the targets.

Basic configuration

Settings in this section cannot be changed after the target group is created.

Choose a target type

- (•) **Instances**
  - Supports load balancing to instances within a specific VPC.
  - Facilitates the use of Amazon EC2 Auto Scaling to manage and scale your EC2 capacity.
- ( ) **IP addresses**
  - Supports load balancing to VPC and on-premises resources.
  - Facilitates routing to multiple IP addresses and network interfaces on the same instance.
  - Offers flexibility with microservice based architectures, simplifying inter-application communication.
  - Supports IPv6 targets, enabling end-to-end IPv6 communication, and IPv4-to-IPv6 NAT.
- ( ) **Lambda function**
  - Facilitates routing to a single Lambda function.
  - Accessible to Application Load Balancers only.
- ( ) **Application Load Balancer**
  - Offers the flexibility for a Network Load Balancer to accept and route TCP requests within a specific VPC.
  - Facilitates using static IP addresses and PrivateLink with an Application Load Balancer.

Target group name

ELBHAZone-tg-group

A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.

Protocol: HTTP : Port: 80

VPC

Select the VPC with the instances that you want to include in the target group.

assignment-3-vpc
vpc-0c80be5bbbb89c469
IPv4: 10.0.0.0/16

Protocol version

- (•) **HTTP1**
  Send requests to targets using HTTP/1.1. Supported when the request protocol is HTTP/1.1 or HTTP/2.
- ( ) **HTTP2**
  Send requests to targets using HTTP/2. Supported when the request protocol is HTTP/2 or gRPC, but gRPC-specific features are not available.
- ( ) **gRPC**
  Send requests to targets using gRPC. Supported when the request protocol is gRPC.

Health checks

The associated load balancer periodically sends requests, per the settings below, to the registered targets to test their status.

Health check protocol: HTTP

Health check path

Use the default path of "/" to ping the root, or specify a custom path if preferred.

/

Up to 1024 characters allowed.

▶ **Advanced health check settings**

Attributes

Certain default attributes will be applied to your target group. You can view and edit them after creating the target group.

▶ Tags - *optional*

Consider adding tags to your target group. Tags enable you to categorize your AWS resources so you can more easily manage them.

Cancel Next

3. create a DB instance in a VPC:

Assuming already we have created VPC

1. We will create a DB subnet group
2. Create a VPC security group
3. Create a DB instance in the VPC

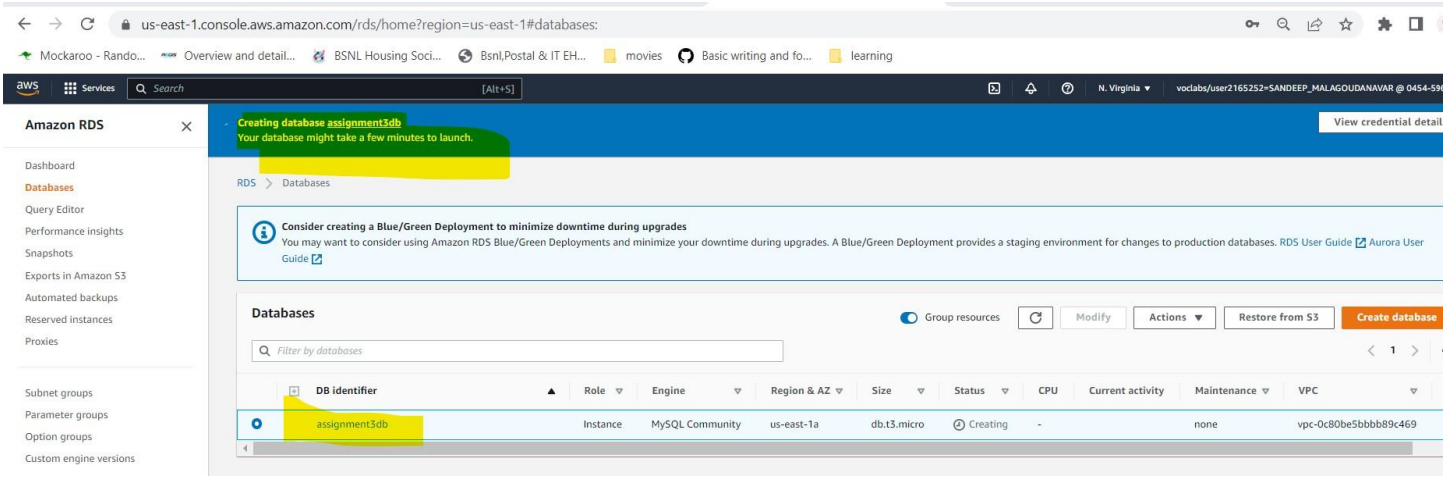

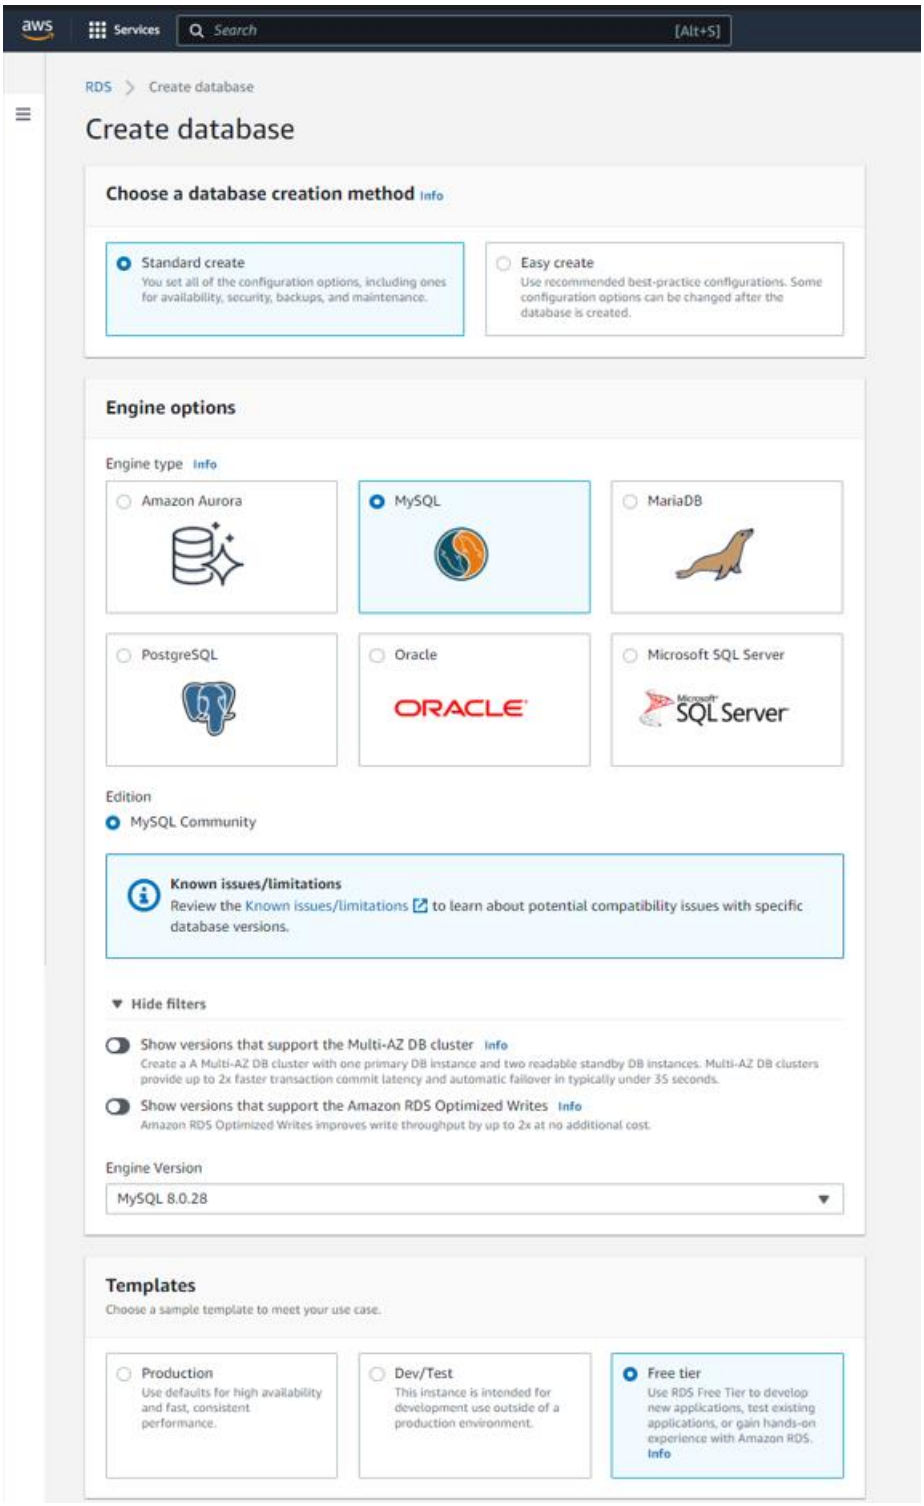
aws
Services
Search
[Alt+S]
RDS
Create database
Create database
Choose a database creation method Info
Standard create
You set all of the configuration options, including ones for availability, security, backups, and maintenance.
Easy create
Use recommended best-practice configurations. Some configuration options can be changed after the database is created.
Engine options
Engine type Info
Amazon Aurora
MySQL
MariaDB
PostgreSQL
Oracle
Microsoft SQL Server
Edition
MySQL Community
Known issues/limitations
Review the Known issues/limitations to learn about potential compatibility issues with specific database versions.
Hide filters
Show versions that support the Multi-AZ DB cluster Info
Create a A Multi-AZ DB cluster with one primary DB instance and two readable standby DB instances. Multi-AZ DB clusters provide up to 2x faster transaction commit latency and automatic failover in typically under 35 seconds.
Show versions that support the Amazon RDS Optimized Writes Info
Amazon RDS Optimized Writes improves write throughput by up to 2x at no additional cost.
Engine Version
MySQL 8.0.28
Templates
Choose a sample template to meet your use case.
Production
Use defaults for high availability and fast, consistent performance.
Dev/Test
This instance is intended for development use outside of a production environment.
Free tier
Use RDS Free Tier to develop new applications, test existing applications, or gain hands-on experience with Amazon RDS.
Info

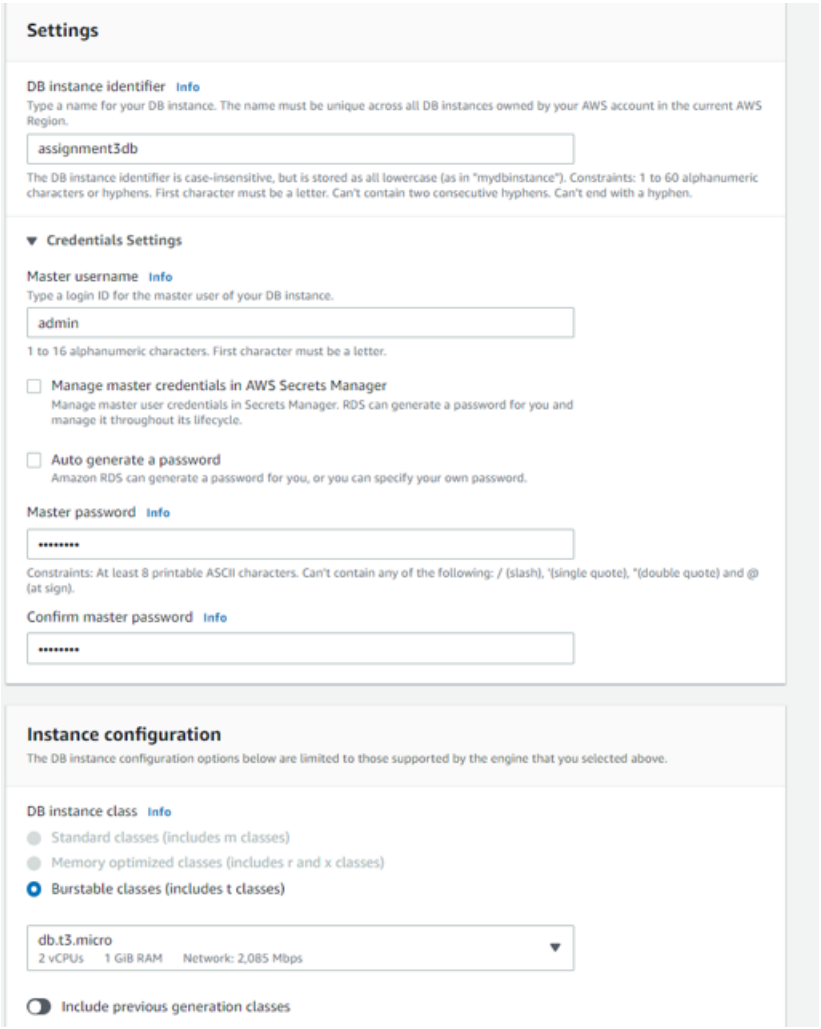

Note : Public access is set to NO. So that we can hide the DB instance from the public access.

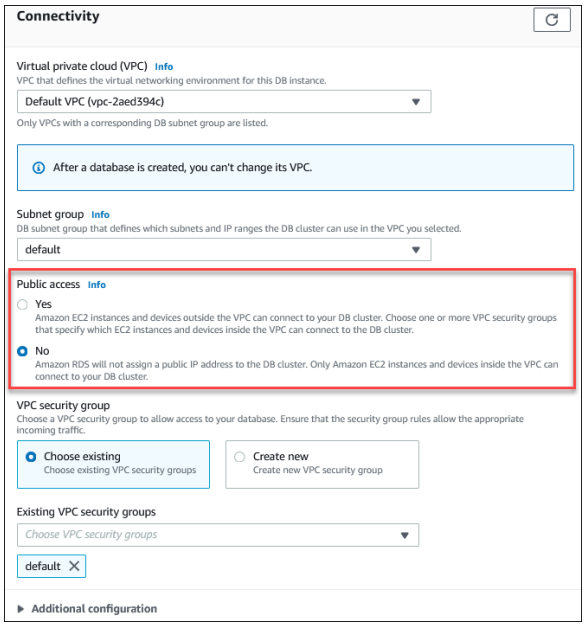

4. create a DB subnet group

1. Open the Amazon RDS console at https://us-east-1.console.aws.amazon.com/rds/home?region=us-east-1
2. choose **Subnet groups**.
3. Choose **Create DB Subnet Group**.
4. **Name**, type the name of your DB subnet group.
5. **Description**, type a description for your DB subnet group.
6. **VPC**, choose the default VPC or the VPC that you created.
7. **Add subnets** section, choose the Availability Zones that include the subnets from **Availability Zones**, and then choose the subnets from **Subnets - private**.

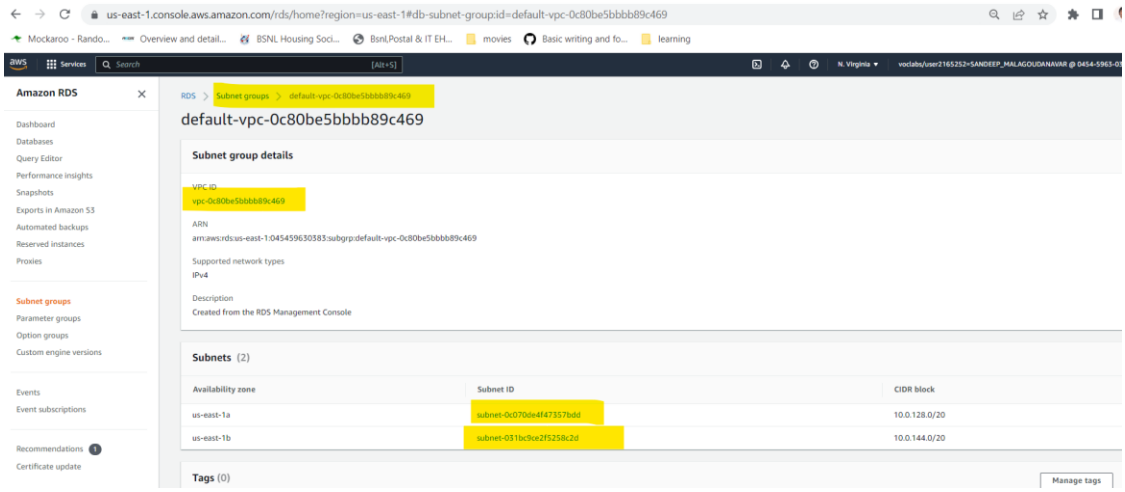

5. Create a load balancer using the AWS Management Console, complete the following tasks.

Configure a target group - Already we create target group.

1. Register targets
2. Configure a load balancer and a listener
3. Test the load balancer

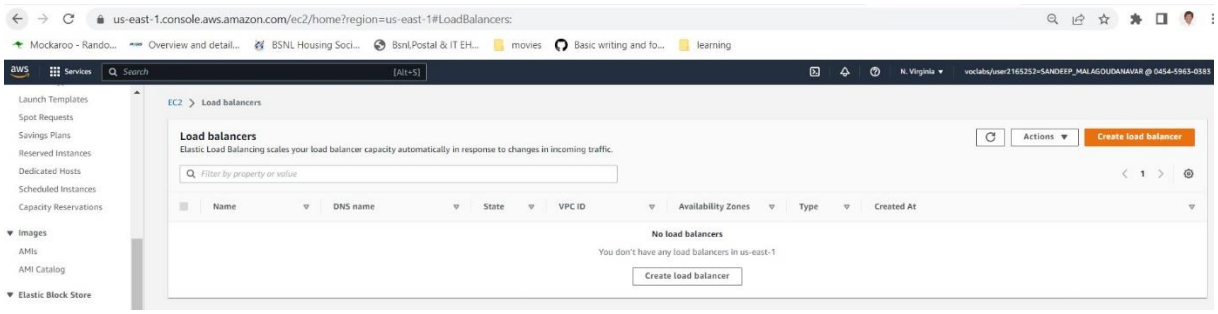

Under **Application Load Balancer**, choose **Create**.

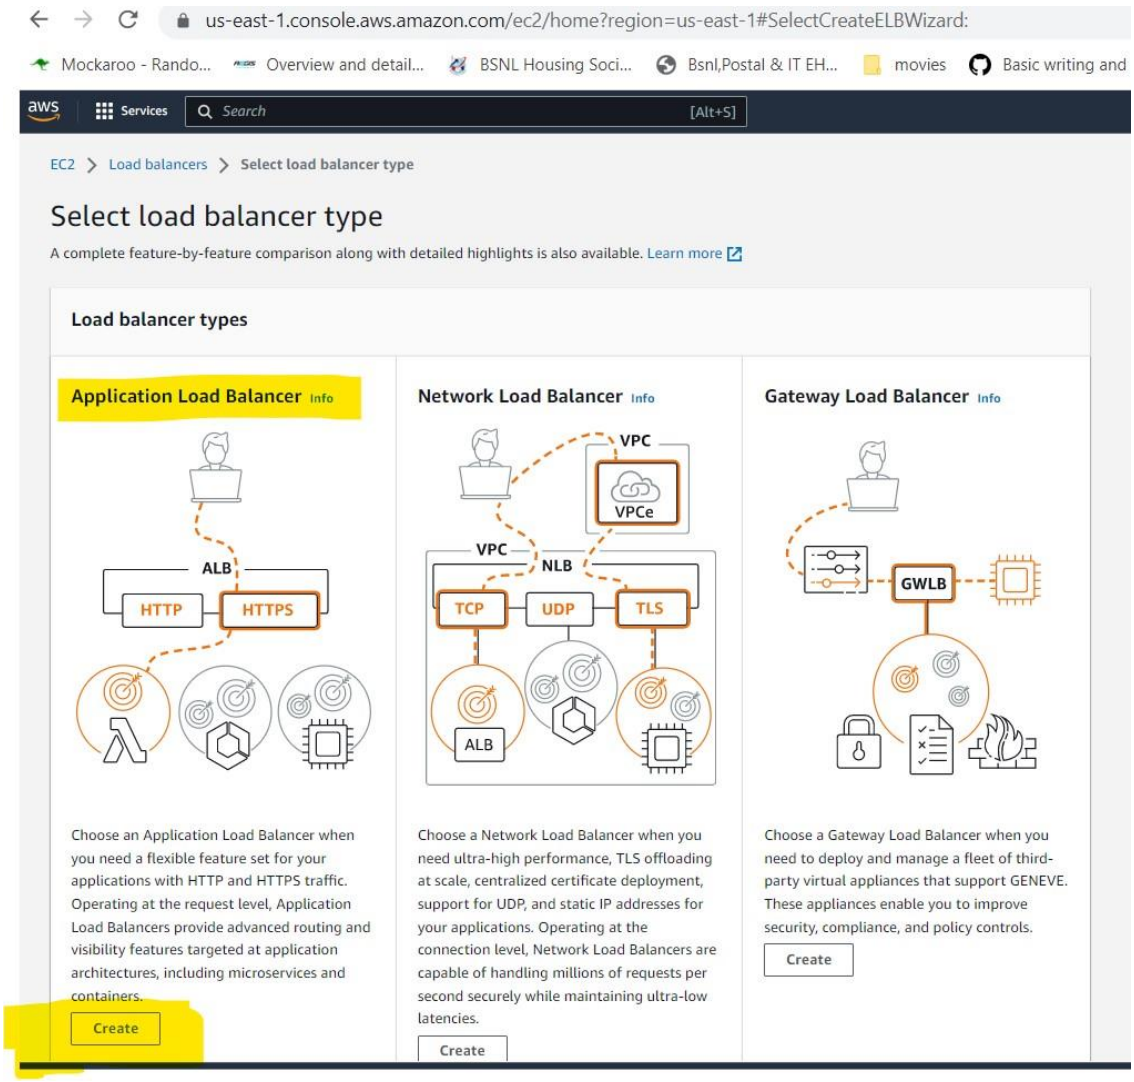

In next steps

- choose **Internet-facing** or **Internal**. An internet-facing load balancer routes requests from clients to targets over the internet. An internal load balancer routes requests to targets using private IP addresses.
- Select an existing security group
- For **Listeners and routing**, the default listener accepts HTTP traffic on port 80. You can keep the default protocol and port, or choose different ones. For **Default action**, choose the target group that you created.

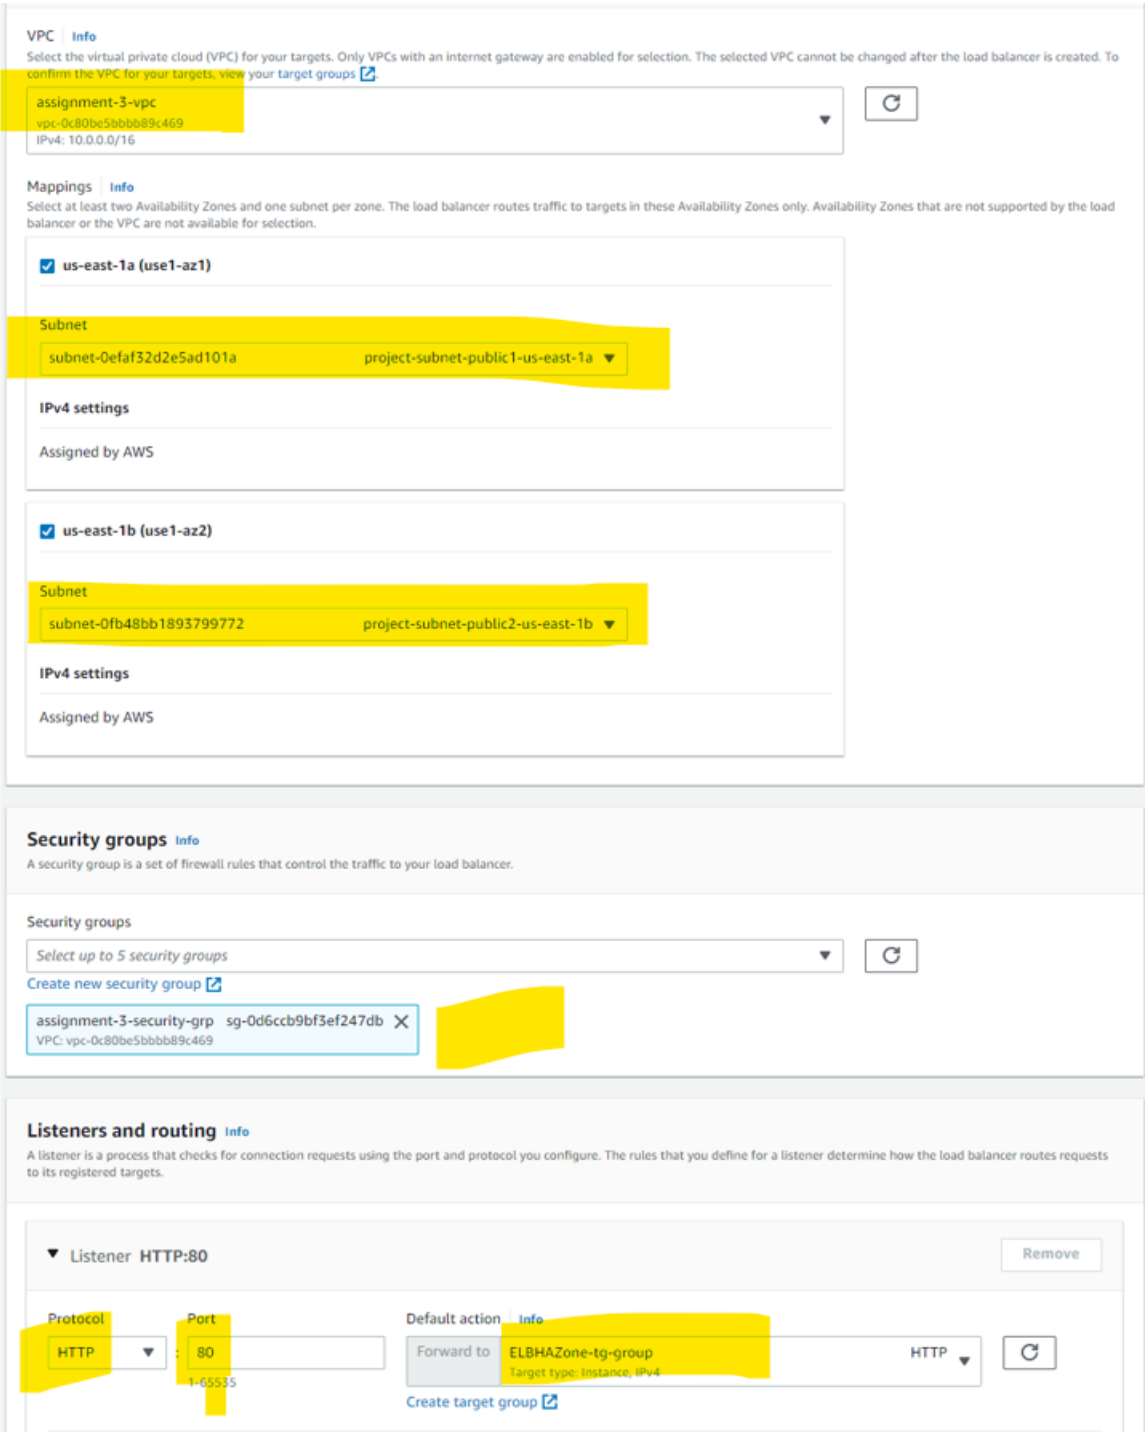
VPC Info
Select the virtual private cloud (VPC) for your targets. Only VPCs with an internet gateway are enabled for selection. The selected VPC cannot be changed after the load balancer is created. To confirm the VPC for your targets, view your target groups
assignment-3-vpc
vpc-0c80be5bbbb89c469
IPv4: 10.0.0.0/16
Mappings Info
Select at least two Availability Zones and one subnet per zone. The load balancer routes traffic to targets in these Availability Zones only. Availability Zones that are not supported by the load balancer or the VPC are not available for selection.
us-east-1a (use1-az1)
Subnet
subnet-0efaf32d2e5ad101a project-subnet-public1-us-east-1a
IPv4 settings
Assigned by AWS
us-east-1b (use1-az2)
Subnet
subnet-0fb48bb1893799772 project-subnet-public2-us-east-1b
IPv4 settings
Assigned by AWS
Security groups Info
A security group is a set of firewall rules that control the traffic to your load balancer.
Security groups
Select up to 5 security groups
Create new security group
assignment-3-security-grp sg-0d6ccb9bf3ef247db
VPC: vpc-0c80be5bbbb89c469
Listeners and routing Info
A listener is a process that checks for connection requests using the port and protocol you configure. The rules that you define for a listener determine how the load balancer routes requests to its registered targets.
Listener HTTP:80
Remove
Protocol
HTTP
Port
80
1-65535
Default action Info
Forward to
ELBHAZone-tg-group
Target type: Instance, IPv4
HTTP
Create target group

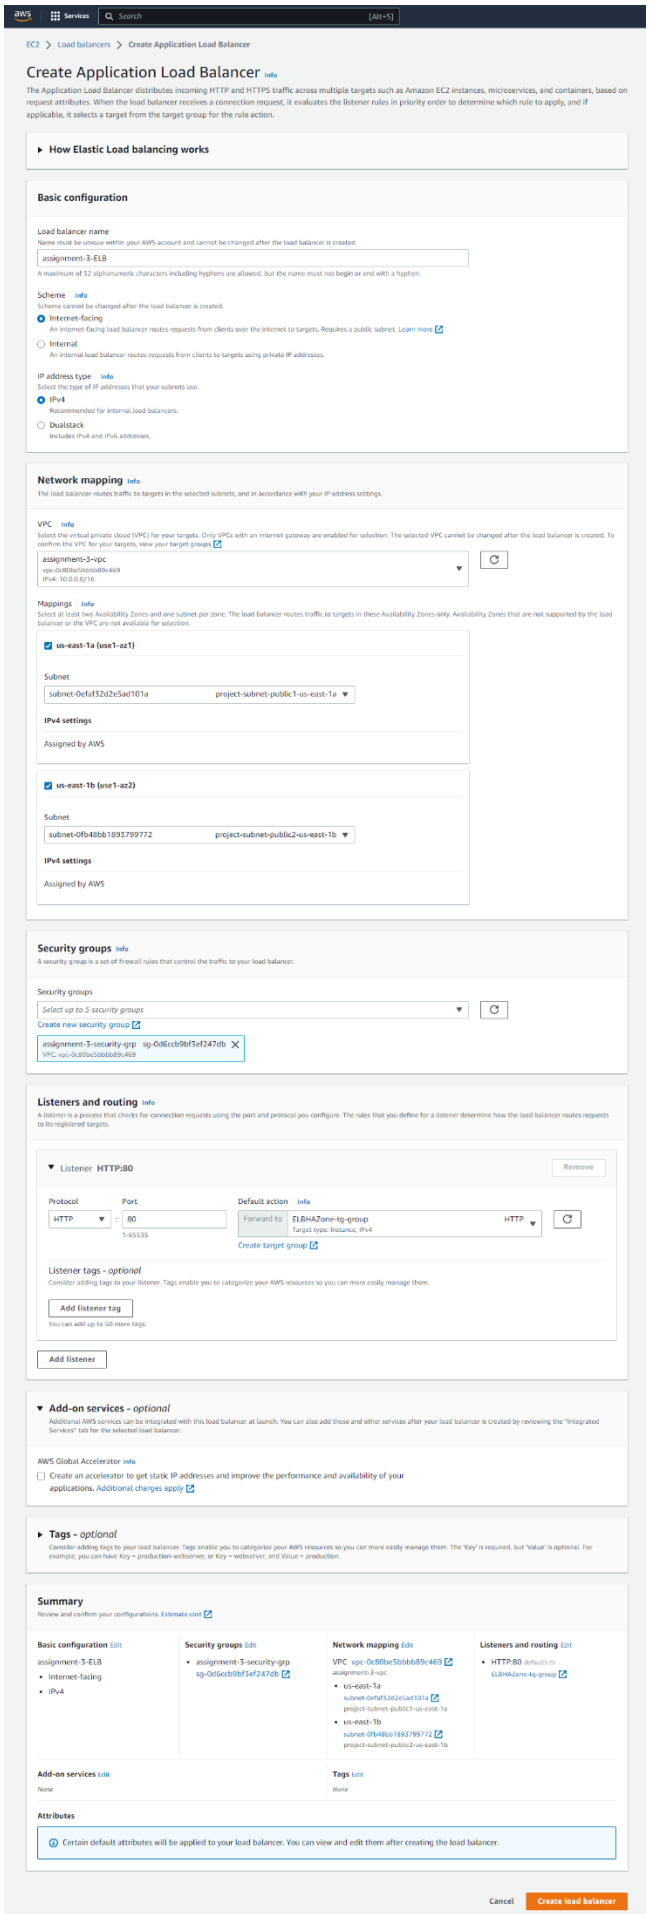
Services
Search
[Alt+S]
EC2 > Load balancers > Create Application Load Balancer
Create Application Load Balancer Info
The Application Load Balancer distributes incoming HTTP and HTTPS traffic across multiple targets such as Amazon EC2 instances, microservices, and containers, based on request attributes. When the load balancer receives a connection request, it evaluates the listener rules in priority order to determine which rule to apply, and if applicable, it selects a target from the target group for the rule action.
How Elastic Load balancing works
Basic configuration
Load balancer name
Name must be unique within your AWS account and cannot be changed after the load balancer is created.
assignment-3-ELB
A maximum of 32 alphanumeric characters including hyphens are allowed, but the name must not begin or end with a hyphen.
Scheme Info
Scheme cannot be changed after the load balancer is created.
Internet-facing
An Internet-facing load balancer routes requests from clients over the Internet to targets. Requires a public subnet. Learn more
Internal
An internal load balancer routes requests from clients to targets using private IP addresses.
IP address type Info
Select the type of IP addresses that your subnets use.
IPv4
Recommended for internal load balancers.
Dualstack
Includes IPv4 and IPv6 addresses.
Network mapping Info
The load balancer routes traffic to targets in the selected subnets, and in accordance with your IP address settings.
VPC Info
Select the virtual private cloud (VPC) for your targets. Only VPCs with an internet gateway are enabled for selection. The selected VPC cannot be changed after the load balancer is created. To confirm the VPC for your targets, view your target groups
assignment-3-vpc
vpc-0c80be5bbbb89c469
IPv4: 10.0.0.0/16
Mappings Info
Select at least two Availability Zones and one subnet per zone. The load balancer routes traffic to targets in these Availability Zones only. Availability Zones that are not supported by the load balancer or the VPC are not available for selection.
us-east-1a (use1-az1)
Subnet
subnet-0efaf32d2e5ad101a
project-subnet-public1-us-east-1a
IPv4 settings
Assigned by AWS
us-east-1b (use1-az2)
Subnet
subnet-0fb48bb1893799772
project-subnet-public2-us-east-1b
IPv4 settings
Assigned by AWS
Security groups Info
A security group is a set of firewall rules that control the traffic to your load balancer.
Security groups
Select up to 5 security groups
Create new security group
assignment-3-security-grp sg-0d6ccb9bf3ef247db
VPC: vpc-0c80be5bbbb89c469
Listeners and routing Info
A listener is a process that checks for connection requests using the port and protocol you configure. The rules that you define for a listener determine how the load balancer routes requests to its registered targets.
Listener HTTP:80
Remove
Protocol
HTTP
Port
80
1-65535
Default action Info
Forward to
ELBHAZone-tg-group
HTTP
Target type: Instance, IPv4
Create target group
Listener tags - optional
Consider adding tags to your listener. Tags enable you to categorize your AWS resources so you can more easily manage them.
Add listener tag
You can add up to 50 more tags.
Add listener
Add-on services - optional
Additional AWS services can be integrated with this load balancer at launch. You can also add these and other services after your load balancer is created by reviewing the "Integrated Services" tab for the selected load balancer.
AWS Global Accelerator Info
Create an accelerator to get static IP addresses and improve the performance and availability of your applications. Additional charges apply
Tags - optional
Consider adding tags to your load balancer. Tags enable you to categorize your AWS resources so you can more easily manage them. The 'Key' is required, but 'Value' is optional. For example, you can have Key = production-webserver, or Key = webserver, and Value = production.
Summary
Review and confirm your configurations. Estimate cost
Basic configuration Edit
assignment-3-ELB
Internet-facing
IPv4
Security groups Edit
assignment-3-security-grp
sg-0d6ccb9bf3ef247db
Network mapping Edit
VPC vpc-0c80be5bbbb89c469
assignment-3-vpc
us-east-1a
subnet-0efaf32d2e5ad101a
project-subnet-public1-us-east-1a
us-east-1b
subnet-0fb48bb1893799772
project-subnet-public2-us-east-1b
Listeners and routing Edit
HTTP:80 defaults to
ELBHAZone-tg-group
Add-on services Edit
None
Tags Edit
None
Attributes
Certain default attributes will be applied to your load balancer. You can view and edit them after creating the load balancer.
Cancel
Create load balancer
Feedback
Looking for language selection? Find it in the new Unified Settings

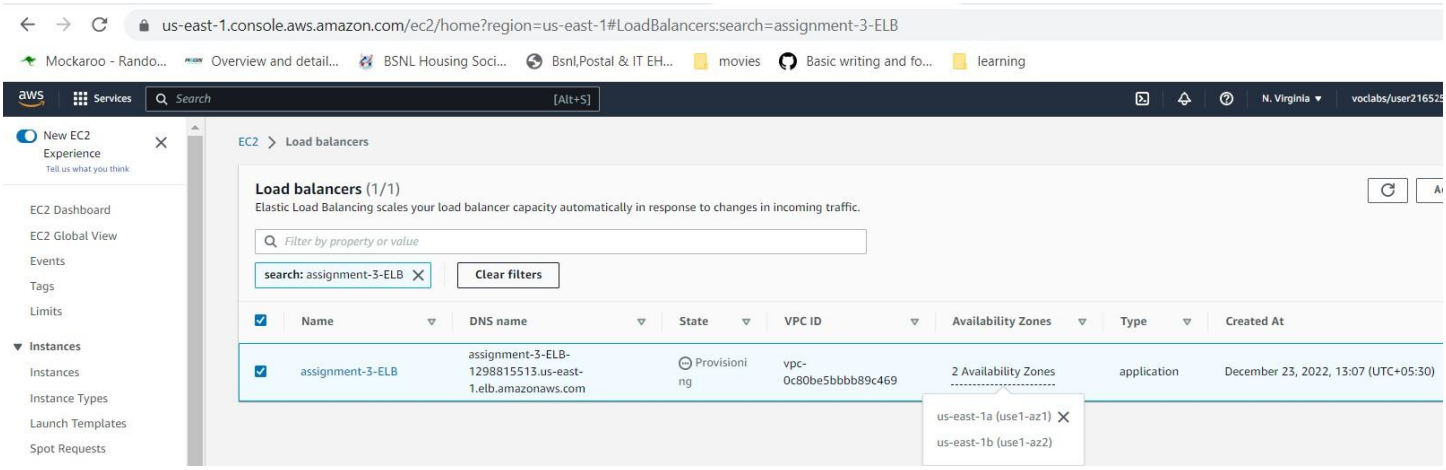

6. Launch configurations

A *launch configuration* is an instance configuration template that an Auto Scaling group uses to launch EC2 instances. When you create a launch configuration, you specify information for the instances. Include the ID of the Amazon Machine Image (AMI), the instance type, a key pair, one or more security groups. You can specify your launch configuration with multiple Auto Scaling groups

- Under Auto Scaling, choose Launch Configurations.
- In the navigation bar, select your AWS Region.
- Choose Create launch configuration and enter a name for your launch configuration.
- For Amazon machine image (AMI), choose an AMI – may be our custom AMI create with webservers.

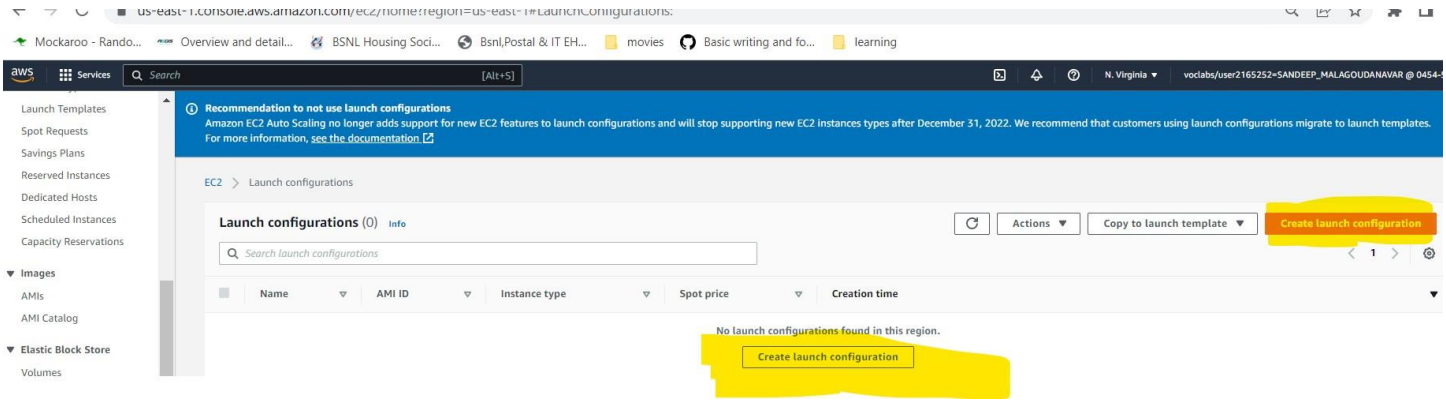

[Alt+S]

reate Auto Scaling group

Choose launch template or configuration Info

Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group. If you currently use launch configurations, you might consider migrating to launch templates.

Name

Auto Scaling group name

Enter a name to identify the group.

assignment-3-auto-scalling-grp

Must be unique to this account in the current Region and no more than 255 characters.

Launch template Info

Switch to launch configuration

Launch template

Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

assignment-3-lunch-cconfig

Create a launch template

Version

Default (1)

Create a launch template version

Description assignment-3-lunch-cconfig-dev	Launch template assignment-3-lunch-cconfig lt-03ca6c11ec0863e45	Instance type t1.micro
AMI ID ami-0574da719dca65348	Security groups -	Request Spot Instances No
Key pair name sandeepm_ec_keypair	Security group IDs -	

Additional details

Storage (volumes) -	Date created Fri Dec 23 2022 13:27:38 GMT+0530 (India Standard Time)	

Cancel Next

[Alt+S]

ate Auto Scaling group

Choose instance launch options Info

Choose the VPC network environment that your instances are launched into, and customize the instance types and purchase options.

Network Info

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-0c80be5bbbb89c469 (assignment-3-vpc)
10.0.0.0/16

Create a VPC

Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets

us-east-1a | subnet-0efaf32d2e5ad101a (project-subnet-public1-us-east-1a)
10.0.0.0/20

us-east-1b | subnet-0fb48bb1893799772 (project-subnet-public2-us-east-1b)
10.0.16.0/20

Create a subnet

Instance type requirements Info

Override launch template

You can keep the same instance attributes or instance type from your launch template, or you can choose to override the launch template by specifying different instance attributes or manually adding instance types.

Launch template	Version	Description
assignment-3-lunch-cconfig lt-03ca6c11ec0863e45	Default	assignment-3-lunch-cconfig-dev

Instance type

t1.micro

Cancel | Previous | Skip to review | Next

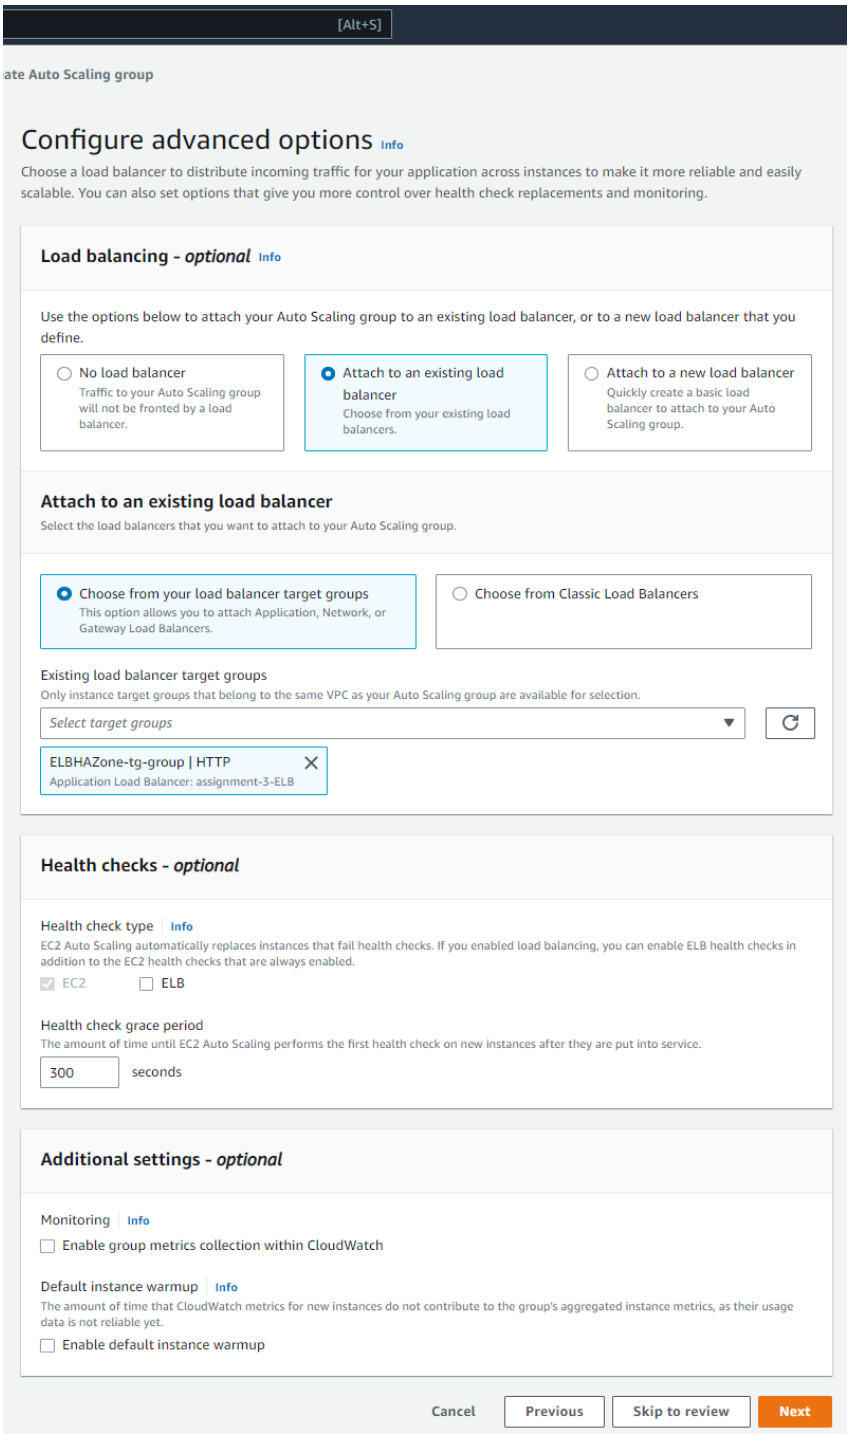
[Alt+S]
ate Auto Scaling group
Configure advanced options Info
Choose a load balancer to distribute incoming traffic for your application across instances to make it more reliable and easily scalable. You can also set options that give you more control over health check replacements and monitoring.
Load balancing - optional Info
Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.
No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.
Attach to an existing load balancer
Choose from your existing load balancers.
Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.
Attach to an existing load balancer
Select the load balancers that you want to attach to your Auto Scaling group.
Choose from your load balancer target groups
This option allows you to attach Application, Network, or Gateway Load Balancers.
Choose from Classic Load Balancers
Existing load balancer target groups
Only instance target groups that belong to the same VPC as your Auto Scaling group are available for selection.
Select target groups
ELBHAZone-tg-group | HTTP
Application Load Balancer: assignment-3-ELB
Health checks - optional
Health check type Info
EC2 Auto Scaling automatically replaces instances that fail health checks. If you enabled load balancing, you can enable ELB health checks in addition to the EC2 health checks that are always enabled.
EC2
ELB
Health check grace period
The amount of time until EC2 Auto Scaling performs the first health check on new instances after they are put into service.
300
seconds
Additional settings - optional
Monitoring Info
Enable group metrics collection within CloudWatch
Default instance warmup Info
The amount of time that CloudWatch metrics for new instances do not contribute to the group's aggregated instance metrics, as their usage data is not reliable yet.
Enable default instance warmup
Cancel
Previous
Skip to review
Next

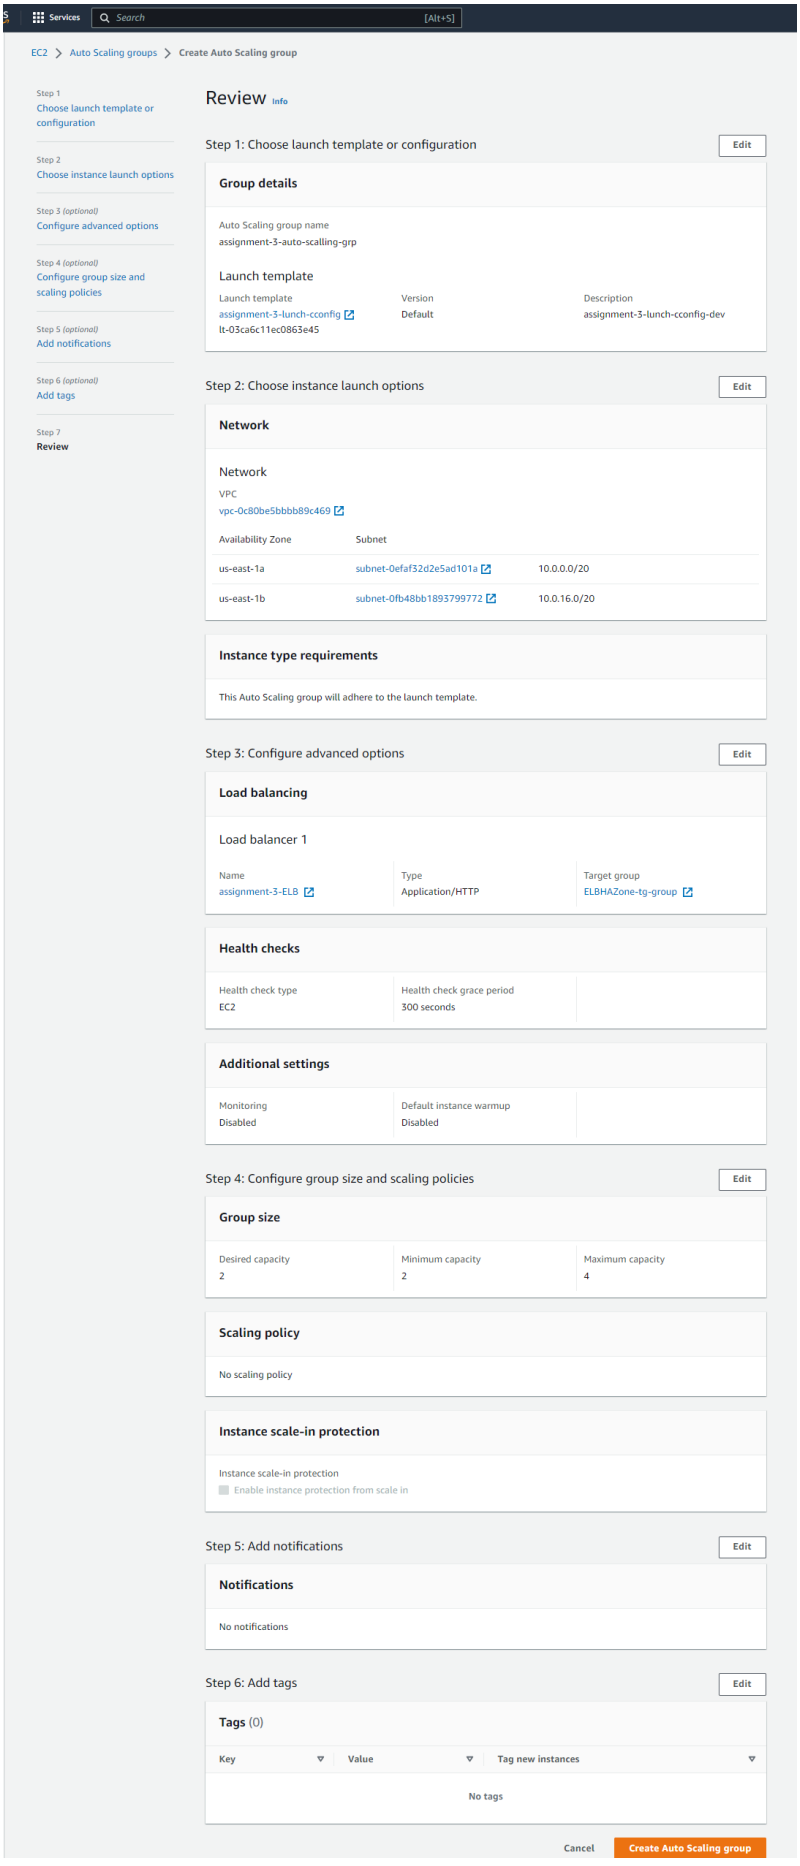
Services
Search
[Alt+S]
EC2 > Auto Scaling groups > Create Auto Scaling group
Step 1
Choose launch template or configuration
Step 2
Choose instance launch options
Step 3 (optional)
Configure advanced options
Step 4 (optional)
Configure group size and scaling policies
Step 5 (optional)
Add notifications
Step 6 (optional)
Add tags
Step 7
Review
Review Info
Step 1: Choose launch template or configuration
Edit
Group details
Auto Scaling group name
assignment-3-auto-scalling-grp
Launch template
Launch template
assignment-3-lunch-cconfig
lt-03ca6c11ec0863e45
Version
Default
Description
assignment-3-lunch-cconfig-dev
Step 2: Choose instance launch options
Edit
Network
Network
VPC
vpc-0c80be5bbbb89c469
Availability Zone
Subnet
us-east-1a
subnet-0efaf32d2e5ad101a
10.0.0.0/20
us-east-1b
subnet-0fb48bb1893799772
10.0.16.0/20
Instance type requirements
This Auto Scaling group will adhere to the launch template.
Step 3: Configure advanced options
Edit
Load balancing
Load balancer 1
Name
assignment-3-ELB
Type
Application/HTTP
Target group
ELBHAZone-tg-group
Health checks
Health check type
EC2
Health check grace period
300 seconds
Additional settings
Monitoring
Disabled
Default instance warmup
Disabled
Step 4: Configure group size and scaling policies
Edit
Group size
Desired capacity
2
Minimum capacity
2
Maximum capacity
4
Scaling policy
No scaling policy
Instance scale-in protection
Instance scale-in protection
Enable instance protection from scale in
Step 5: Add notifications
Edit
Notifications
No notifications
Step 6: Add tags
Edit
Tags (0)
Key
Value
Tag new instances
No tags
Cancel
Create Auto Scaling group